<서울대학교 자연과학대학 생명과학부 학사졸업논문>

Petasearch: Fast, approximate comparison of huge sequence datasets

(페타탐색: 방대한 서열 데이터셋에 대해 빠른 유사성 검색)

이 논문을 이학사 학위논문으로 제출함

2022년 8월

구분: 실험

논문심사 대상자 소속: 자연과학대학 생명과학부

학번: 2017-17232

성명: Minghang Li   (인)

논문지도교수: Martin Steinegger   (인)

연구윤리 준수 서약서

본인 (Minghang Li)은 서울대학교 연구자로 연구를 진행함에 있어 다음 사항을 준수할 것을 서약합니다.

1. 서울대학교 연구윤리 관련 규정 및 지침과 국가 법령 및 정부 지침 그리고 일반적으로 학계에서 인정되는 연구윤리 기준을 준수하여 서울대학교 내에서 연구를 수행할 때 위조 변조 표절 등 학문적 진실성을 훼손하는 연구부정행위 또는 연구부적절행위를 하지 않겠습니다.

2. 인간, 동물 등 연구대상에 대한 국내외 윤리 기준을 준수하도록 하겠습니다.

3. 연구 진행 중 이해상충이 발생할 경우 이를 공개하도록 하겠습니다.

4. 정부 및 본교 지침에 따라 연구노트를 작성하며, 연구 데이터에 대한 관리를 철저히 하도록 하겠습니다.

<div align="center">

2022년 6월 1일

서약자 소속(학과명): 자연과학대학 생명과학부

이름:　Minghang Li　(서명)

</div>

Abstract

The Sequence Read Archive currently holds over 60 petabases and representing a treasure trove for medicine and biotechnology. Bloom-filter and sketching based approaches were proposed to accelerate searches, however they offer only limited sensitivity. We developed Petasearch to enable fast and sensitive searching through huge protein databases. Its algorithm contains three stages: (1) We pre-process the database sequences to extract k-mers, sort and store them in a highly compressed k-mer index. (2) We extract query k-mers, add similar k-mers and find matches between query and database k-mers. To maximize throughput, we exploit the caching and prefetch infrastructure of modern CPUs, advanced Linux IO techniques, and the enormous read bandwidth of NVMe-SSDs. (3) We compute SIMD-accelerated banded Smith-Waterman alignments between sequences of high-scoring k-mer matches. With such design, Petasearch is proved to have great efficiency: it is up to 190 times faster than state-of-the-art algorithms on a 9.3TB benchmark. At much accelerated speeds, Petasearch matches state-of-the-art algorithms on sensitivity down to sequence identities of 60%. On a SCOP25 benchmark we showed that Petasearch's profile search detects sequence homology down to 40% sequence identity. We also showed that Petasearch can be applied in finding novel Cas family proteins and discovering new RNA-dependent RNA polymerase (RdRP) homologs. In conclusion, Petasearch is a tool with huge potential. It will enable fast querying of current and upcoming databases and bring bioinformatic researches to a larger scale.

....

Contents

1. Introduction

1.1 Sequence Databases

Next generation sequencing (NGS) technologies have revolutionized the way we collect and analyze biological data. Thanks to NGS, the cost of sequencing has dropped drastically and continued to decrease with more new technologies developed. Accompanying this change is the explosive growth of the amount of sequencing data and the size of sequence databases. The Sequence Read Archive (SRA) is one of the most popular and widely used sequence databases that store both private and public sequence reads and provide access in various foramts including the commonly used `FASTQ` file format. Its size has grown exponentially since 2008 and currently reached more than 60 petabytes large. The growth in size of Sequence Read Archive is visualized in Figure 1.1.

1.2 State-of-the-art Algorithms for Sequence Searches

1.2.1 DIAMOND

1.2.2 MMseqs2

1.2.3 BIGSI

1.3 Prototype of Petasearch Algorithm

1.4 Motivation and Contribution of the Thesis

The search of homologs in large sequence databases requires a fast yet sensitive enough algorithm specially designed for petabyte-scale analysis. The state-of-the-art searching algorithms failed to satisfy this need. The prototype of `Petasearch`, despite its idea proven to be promising, has not reached its peak speed efficiency and is rather heavy in disk consumption. Limited by the current design, its searching sensitivity is also less desirable for homologs with sequence similarity less than 40%. To tackle these problems and make `Petasearch` more available to the public, we revised the design of the core data structures of `Petasearch` and added the profile-search functionality. The main contribution of this thesis is the major improvement of the `Petasearch` algorithm in speed, space consumption

and sensitivity.

In Chapter 2, we will continue with describing the further development and optimization of `Petasearch`. We will also describe the design of the benchmarks in Chapter 2. In Chapter 3, we will first show the improvements in efficiency and effectiveness of the forementioned optimizations. Afterwards, we will show a thorough comparison of the performance of the `Petasearch` algorithm with the state-of-the-art algorithms. In Chapter 4, we will discuss the potential application of `Petasearch` and show two examples of its usage.

Figure 1.1: The exponential growth of the Sequence Read Archive from 2008 to 2022. The total amount of sequence data (unit in bases) and publicly available data are visualized in pink and dark red respectively.

2. Materials & Methods

2.1 Space Optimization

The core idea of `Petasearch` is actually sacrificing space for fast computation. However, the resources are not limitless. Thus, we would like to keep the cost of space as low as possible while keeping the searching speed high. In this section, we will discuss the space optimization techniques utilized to improve `Petasearch` prototype.

2.1.1 Diff-index Compression

As is described in chapter 2, the diff-index created in the k-mer extraction step will store multiple `USHRT_MAX` as long as the difference is larger than `USHRT_MAX`. This will make any k-mer difference larger than $4 \times$ `USHRT_MAX` $= 262140$ require a larger space to store than the original `unsigned long` representation. This situation is not uncommon especially when k is large.

Also, in the prototypical implementation, the ID of the source sequence will also be stored multiple times in the ID table. This redundancy is both unnecessary and troublesome. It will increase the size of `Petasearch` data structures even more than the repeated `USHRT_MAX` since the IDs are stored as `unsigned long` (64-bit integers).

Figure 2.1 showed the space consumption of the diff-index created in the k-mer extraction step when $k = 11$. Without optimization, the diff-index (k-mer table) and its corresponding ID table will take up 17 GB of space for a merely 1GB-sized database.

To optimize the size of the diff-index, we devised the bit-squeezing technique to compress the difference between two adjacent k-mers: For any 64-bit k-mer difference, we continuously fetch 15 bits into a write buffer starting from the least significant bits. We stop the retrieval until we encounter a zero chunk (15 bits of zeros).

To enable the correct decoding of the diff-index during the next phase, the sign bit of the last element in th write is set to `1` to indicate the end of encoding. Afterwards, we write all the elements in the write buffer to the diff-index. An example encoding process for differecne of value 2039432531946 is shown in Figure 2.2. For ID table, the optimization is simple: we store the ID of the source sequence only once instead of repeatedly.

Using the bit-squeezing method, it is possible to obtain a maximum of five chunks, making the final space consumption larger than the size of a `unsigned long` integer. However, such situation only happens when the difference is larger than `1UL << 59 =` 576460752303423488, which is extremely rare.

While decoding the compressed diff-index in the process of double-index search, we will reverse the bit-squeezing process through repeatedly retrieving 15 bits from the diff-index table until we encounter the chunk with the sign bit set to `1`. The decoded difference value will be add to the current k-mer. Moreover, since we do not store redundant IDs, the ID pointer will not be incremented until the end of k-mer decoding. Algorithm 1 showed the simplified pseudocode for k-mer decoding.

Algorithm 1 Pseudocode for the k-mer decoding process
---
procedure DecodeKmer(*currentKmer, currentTargetKmerPtr, currentTargetIDPtr*)

$\quad$ *currentDiffIndex* ← 0

$\quad$ while ∗*currentTargetKmerPtr* > 0 do $\qquad\qquad$ ▷ This means the sign bit is not 1.

$\quad\quad$ *currentDiffIndex* ← Get15Bits(∗*currTargetKmerPtr*)

$\quad\quad$ *currentDiffIndex* ← *currentDiffIndex* << 15

$\quad\quad$ Next(*currTargetKmerPtr*)

$\quad$ end while

$\quad$ *currentDiffIndex* ← Get15Bits(∗*currTargetKmerPtr*)

$\quad$ *currentKmer* ← *currentKmer* + *currentDiffIndex*

$\quad$ Next(*currTargetKmerPtr*)

$\quad$ *currentTargetIDPtr* ← Next(*currentTargetIDPtr*) return *currentKmer*

end procedure

2.1.2 Protein Sequence Compression

For terabyte-size databases, the sequences themselves are also space consuming. To further reduce the size of the databases, we developed the `ASCII`-sqeezing technique.

Protein sequences are represented by a limited subset of `ASCII` characters, which are encoded by a single byte. However, as is shown in Figure 2.3, we only need 5 bits to represent all the amino acids. Therefore, we can squeeze every three amino acids into one 16-bit `short`. Similar to the bit-squeezing technique described in Section 2.1.1, we also use the sign bit to indicated the end of the compressed protein sequence. Figure 2.4 showed an example compression process for glutathione (GSH). The `ASCII`-squeezing technique is expected to produce a sqeuence database about 85% of the original size.

Figure 2.1: Visualizaiton of k-mer table and ID table sizes when $k = 11$. The database is the **UniProtKB/Swiss-Prot** database obtained through `mmseqs databases UniProtKB/ Swiss-Prot swissprot tmp` command. Without optimization, the sizes of `petasearch` data structures are 6.46 times and 12.92 times larger than the sequence database.
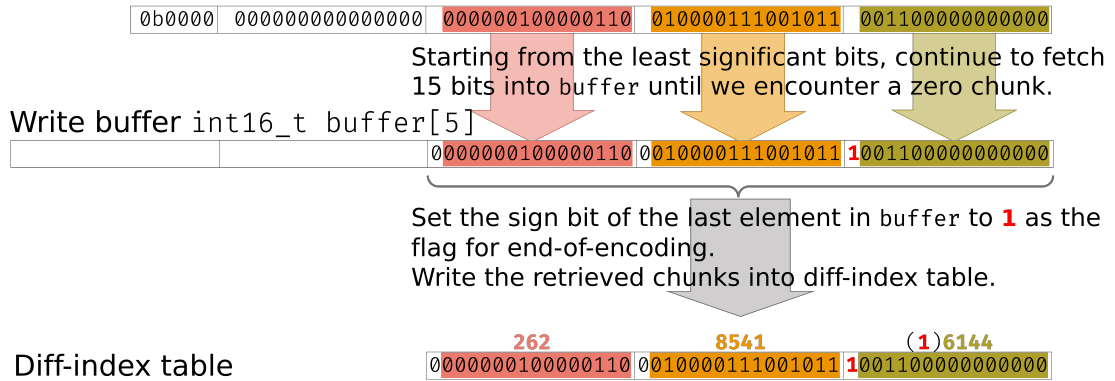

Figure 2.2: The example decoding process of difference index 2039432531946. We first retrieve 15-bit chunks starting from the least significant bits and store them into a write buffer in the reverse order until we encounter a zero chunk. For 2039432531946, its highest non-zero bit is 39, which means that we need three 16-bit `short` to store it.

2.2 Speed Optimization

Speed is the first and foremost concern of `Petasearch`. The speed of `Petaserch` prototype is already fast, but did not make it stand out too much from its competitors. In this section, we will introduce several techniques to further boost the speed of `Petasearch`.

2.2.1 IO Performance Optimization

In the prototypical `Petasearch` implementation, `mmap` was selected for reading `Petasearch` data structures stored on NVMe SSDs. However, `mmap` does not scale well with the increase in threads (Papagiannis et al., 2020) and thus cannot saturate the full throughput of NVMe SSDs. To find the IO tool with the best performance, we conducted a benchmarking study on the performance of various IO tools using FIO benchmark software (Axboe, 2022). For synced IO tools, we benchmarked `pread` using different flags and `mmap`. For async IO tools, we benchmarked `libaio` and `posix_aio`.

The benchmarking results are visualized in Figure 2.5. It is clear that `libaio` performs the best. I is able to saturate the full 3.5 GB/s linear read bandwidth of NVMe SSDs. The other two tools, `posix_aio` and `pread` with `O_DIRECT` (`ioengine = psync` in FIO) have roughly the same performance, with bandwidth around 3.3 GB/s. Unfortunately, `mmap` has the worst performance, with only about 1.5 GB/s at maximum. The performance even fell to 0.25 GB/s when it scaled to 20 threads. Since `mmap` is a synced IO module, adopting another synced IO module will require almost no change in control logic. Considering both the performance and difficulty of refactoring, we chose to use `pread` with `O_DIRECT` in place of `mmap`.

The implementation of `pread` with `O_DIRECT` is rather simple: we simply create a read buffer according to the currently available memory size, open the k-mer diff-index file with `O_DIRECT` flag, and then read in parallel using `pread` continuously untill `EOF` (end of file).

2.2.2 Simplified Database Index

In `Pteasearch` prototype, the sequence database format is the same as that of `MMseqs`, which has many functions uses a complexed index structure. Such complexity is unnecessary for `Petasearch`. Hence, we simplified the index, only preservingthe offset of the corresponding entry in the sequence database. It is expected to reduce the IO time.

A	01000001	J	01001010	S	01010011
B	01000010	K	01001011	T	01010100
C	01000011	L	01001100	U	01010101
D	01000100	M	01001101	V	01010110
E	01000101	N	01001110	W	01010111
F	01000110	O	01001111	X	01011000
G	01000111	P	01010000	Y	01011001
H	01001000	Q	01010001	Z	01011010
I	01001001	R	01010010		

Figure 2.3: Part of the ASCII table, showing the bit representation of `A` to `Z` with the last 5 bits highlighted. It can be clearly seen that for `A` to `Z` in the English alphabet, we can represent them using only 5 bits instead of a whole byte (8 bits).
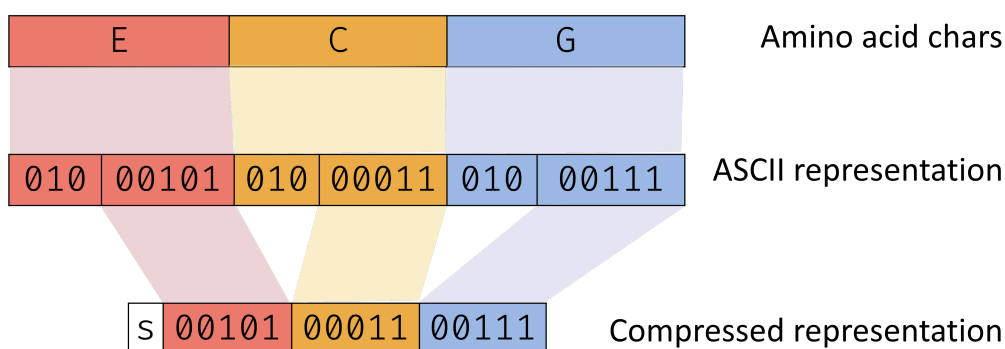

Figure 2.4: The example compression of short peptide glutathione (GSH). GSH consists of only three amino acids: glutamate (E), cysteine (C), and glycine (G). We simply fetch the least significant 5 bits of each amino acid `char` and store them into a single 16-bit `short`.


Figure 2.5: Benchmarking results of various IO tools using FIO benchmark software. The benchmark setting is imitating a parallel read from 20 NVMe SSDs. A total of 20 threads were created, each one responsible of reading a 50 GB file stored on one NVMe SSD. The bandwidth is the average reading bandwidth per SSD. The IO engine `psync` is equivalent to opening a file handle with `O_DIRECT` and reading from the handle using `pread`.

2.2.3   Fast Third-Party Libraries

We integrated several fast thrid-party libraries to replace the slow implementaions in `Petaserch` prototype. The fast parallel in-place sorting algorithm IPS^4o (Axtmann et al., 2017) was integrated to replace the slow `std::sort`. The banded Smith-Waterman-Gotoh aligner `block-aligner` (Liu and Steinegger, 2021) was used to allow fast pairwise alignment in the third phase.

2.3   Sensitivity Improvement

The k-mer matching mechanism limits `Petasearch`'s ability of finding homologs with low sequence identity. To improve `Petasearch`'s performance at lower sequence identity, we made `Petasearch` able to perform profile search by allowing profile databases as inputs.

Profile search is using a "sequence profile" generated from multiple sequence alignment (MSA) results as the querying input (Steinegger et al., 2019). The profile Hidden Markov Model (HMM) provides position-specific aminoaid indel and substitution penalties (Steinegger et al., 2019), which significantly increase the searching sensitivity. The most sensitive searching tools such as `HMMER` (Eddy, 2009, Eddy, 2011), `HHblits` (Remmert et al., 2012) and `HH-suite3` (Steinegger et al., 2019) all use the profile search mechanism. Thus, enabling `Petaserch` to perform profile search is expected to improve its sensitivity at low sequence identity.

3.   Results

4.   Discussion

5. References

Axboe, J. (2022). Flexible I/O Tester. https://github.com/axboe/fio

Axtmann, M., Witt, S., Ferizovic, D., & Sanders, P. (2017). In-place parallel super scalar samplesort (ips 4 o). arXiv preprint arXiv:1705.02257.

Eddy, S. R. (2009). A new generation of homology search tools based on probabilistic inference. In Genome informatics 2009: Genome informatics series vol. 23 (pp. 205–211). World Scientific.

Eddy, S. R. (2011). Accelerated profile hmm searches. PLoS computational biology, 7(10), e1002195.

Liu, D., & Steinegger, M. (2021). Block aligner: Fast and flexible pairwise sequence alignment with simd-accelerated adaptive blocks. bioRxiv.

Papagiannis, A., Xanthakis, G., Saloustros, G., Marazakis, M., & Bilas, A. (2020). Optimizing memory-mapped i/o for fast storage devices. 2020 USENIX Annual Technical Conference (USENIX ATC 20), 813–827.

Remmert, M., Biegert, A., Hauser, A., & Söding, J. (2012). Hhblits: Lightning-fast iterative protein sequence searching by hmm-hmm alignment. Nature methods, 9(2), 173–175.

Steinegger, M., Meier, M., Mirdita, M., Vöhringer, H., Haunsberger, S. J., & Söding, J. (2019). Hh-suite3 for fast remote homology detection and deep protein annotation. BMC bioinformatics, 20(1), 1–15.